

APPLIED RESEARCH PROJECT · CAPSTONE PRESENTATION

A Decision-Oriented Conversational Shopping Assistant

Preference Elicitation · Explainable Recommendations · Lifecycle Support

Chenghui Tan

California State University, East Bay · April 2026

Information Retrieval $\neq$ Decision Support



Cognitive Overload

Users must mentally juggle all constraints and trade-offs themselves



Vague, Evolving Needs

Users rarely know exactly what they want when they start



Preferences Get Overwritten

Add a new constraint — your earlier choices are erased



Missing Critical Info

Delivery time, promos, and availability hidden or absent

The Decision Burden Falls on the User

*Existing tools assume shopping is a search problem.
In reality, users face a decision problem —
and no platform is built to help them solve it.*

- Decision fatigue increases with more options
- Users guess, settle, or abandon purchases
- Post-purchase regret when trade-offs aren't clear
- Trust erodes when recommendations have no rationale

→ We built a system that addresses the decision — not just the search.

Existing Platforms — Where They Fall Short

Google
Shopping
Search & filter

Amazon
Rufus
AI chat in Amazon

ChatGPT
Shopping
Natural language AI

Perplexity
AI answer engine

KEY GAPS ACROSS THE LANDSCAPE

✗ Overwrites preferences mid-conversation

✗ Assumes you already know what you want

✗ Missing delivery, promos & availability

✗ Text heavy — low information density

✗ No preference continuity across turns

✗ Multiple steps to reach product page

✗ Separates reasoning from shopping flow

✗ Opaque ranking — no explanation why

✗ Not built for first-time / cold-start users


→ The common thread: no platform treats shopping as a decision problem.

Search-first tools
→ retrieve results, not decisions

AI chatbots
→ converse, but don't structure trade-offs

All surveyed platforms
→ no preference continuity across turns

Your Personal AI Shopping Assistant




Understand You

Frustration → Category

Decision-Relevant Q&A

ecosystem · placement · privacy · size



Remembers You

Preference Continuity

Visible Diff on Refine

"leaks" → leak-proof · "heavy" → lightweight



Clear Decisions

3 Curated Picks

Why-this-fits · Compare

Best Fit · Budget Pick · Stretch Pick




Lifecycle Layer (prototype) — Realistic post-purchase support

Five-Layer Decision Support Framework




End-to-end: vague need → structured decision → ranked products → explained trade-offs → lifecycle value

Components & How They Communicate



The Deterministic Core



The Supporting Stack — What Feeds the Decision Engine

ETL + Provenance Flags

Playwright scrapers • Amazon • Target
Anti-bot stealth • normalize • dedup
Each field flagged: scraped vs estimated
vs synthesised — no silent fabrication

100 products • 3 categories

Attribute Extraction

screen_inches • ecosystems • has_camera
mounting • capacity_oz • bottle_material
drinking_style • organiser_material
Pulled from titles at load time

feeds ranker, picks, compare view

Claude Opus 4.7 — 3 Call Sites

- ① Parse free-text input
 - ② Write personal product copy
 - ③ Parse refine intent
- Every call has a deterministic fallback

Never on the ranking path

React Frontend — 5 Stages

- ① Need Discovery
- ② Clarification Q&A
- ③ 3 Curated Picks
- ④ Why This Fits (rule-grounded checklist)
- ⑤ Lifecycle (prototype value-extension)

+ Compare overlay • Save-for-later • Diff

Capability Validation — What I Checked

Decision Transparency

How users can see and audit every recommendation



Match / Miss / ? checklist

Every chip the user picked is shown as ✓ match • ✗ miss • ? unknown.



Trade-offs visible

Yellow trade-off card on each pick names the cost, not just the win.



Reasoning inspectable

Each reason maps to a rule the deterministic ranker actually fired.

Capability Check

- Main gaps addressed
- Preference continuity
- Decision modeling
- Comparison + lifecycle

Feedback & Review

- Small-scale user feedback
- Professor feedback
- Industry mentor input
- Implementation review

Result: vague needs → structured preferences → explainable recommendations

Honest Assessment • What This Proves • Next Steps

Limitations

- ⚠ Small dataset — 100 products, 3 categories only
- ⚠ Small-scale user testing — family, friends, classmates only
- ⚠ 72/100 delivery values are estimated (flagged in data)
- ⚠ Rule-based ranking — weights inferred, not learned

Conclusions & Impact

- ✓ Decision support feasible end-to-end without an LLM key
- ✓ Explanations grounded in the same rules the ranker used
- ✓ Visible preference continuity — every refine shows a diff
- ✓ Lifecycle layer is a prototype value-extension concept
- ✓ AI tools disclosed: Claude Opus 4.7 • Codex 5.5 review

Future Work

- Formal user study (SUS, task completion time)
- Live retailer API for real-time pricing
- ML-based ranking from interaction data
- Expand to more product categories
- Mobile-first UI redesign

Thank You

Questions & Discussion

Author:	Chenghui Tan
Institution:	California State University, East Bay (CSUEB)
Project:	Decision-Oriented Conversational Shopping Assistant
Runtime:	React · FastAPI · Claude Opus 4.7 · Python · Playwright (ETL)
Built with:	Claude Code (Opus 4.7) · Codex GPT-5.5 (independent validation)